

# Continuous\_Control

October 2, 2020

## 1 Continuous Control

### 1.0.1 1. Start the Environment

```
[1]: from unityagents import UnityEnvironment
import numpy as np
```

Next, we will start the environment! *Before running the code cell below*, change the `file_name` parameter to match the location of the Unity environment that you downloaded.

- **Mac**: "path/to/Reacher.app"
- **Windows (x86)**: "path/to/Reacher\_Windows\_x86/Reacher.exe"
- **Windows (x86\_64)**: "path/to/Reacher\_Windows\_x86\_64/Reacher.exe"
- **Linux (x86)**: "path/to/Reacher\_Linux/Reacher.x86"
- **Linux (x86\_64)**: "path/to/Reacher\_Linux/Reacher.x86\_64"
- **Linux (x86, headless)**: "path/to/Reacher\_Linux\_NoVis/Reacher.x86"
- **Linux (x86\_64, headless)**: "path/to/Reacher\_Linux\_NoVis/Reacher.x86\_64"

For instance, if you are using a Mac, then you downloaded `Reacher.app`. If this file is in the same folder as the notebook, then the line below should appear as follows:

```
env = UnityEnvironment(file_name="Reacher.app")
```

I have already downloaded the unity enviroment and place it same location as this file, once donde that you can access trough:

```
[2]: env = UnityEnvironment(file_name='Reacher_Linux/Reacher.x86')
```

```
INFO:unityagents:
'Academy' started successfully!
Unity Academy name: Academy
    Number of Brains: 1
    Number of External Brains : 1
    Lesson number : 0
    Reset Parameters :
        goal_speed -> 1.0
        goal_size -> 5.0
Unity brain name: ReacherBrain
    Number of Visual Observations (per agent): 0
    Vector Observation space type: continuous
```

```
Vector Observation space size (per agent): 33
Number of stacked Vector Observation: 1
Vector Action space type: continuous
Vector Action space size (per agent): 4
Vector Action descriptions: , , ,
```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```
[3]: # get the default brain
brain_name = env.brain_names[0]
brain = env.brains[brain_name]
```

## 1.0.2 2. Examine the State and Action Spaces

In this environment, a double-jointed arm can move to target locations. A reward of +0.1 is provided for each step that the agent's hand is in the goal location. Thus, the goal of your agent is to maintain its position at the target location for as many time steps as possible.

The observation space consists of 33 variables corresponding to position, rotation, velocity, and angular velocities of the arm. Each action is a vector with four numbers, corresponding to torque applicable to two joints. Every entry in the action vector must be a number between -1 and 1.

Run the code cell below to print some information about the environment.

```
[6]: # reset the environment
env_info = env.reset(train_mode=True)[brain_name]

# number of agents
num_agents = len(env_info.agents)
print('Number of agents:', num_agents)

# size of each action
action_size = brain.vector_action_space_size
print('Size of each action:', action_size)

# examine the state space
states = env_info.vector_observations
state_size = states.shape[1]
print('There are {} agents. Each observes a state with length: {}'.
      ↪format(states.shape[0], state_size))
print('The state for the first agent looks like:', states[0])
```

```
Number of agents: 1
Size of each action: 4
There are 1 agents. Each observes a state with length: 33
The state for the first agent looks like: [ 0.00000000e+00 -4.00000000e+00
```

```

0.00000000e+00  1.00000000e+00
-0.00000000e+00 -0.00000000e+00 -4.37113883e-08  0.00000000e+00
 0.00000000e+00  0.00000000e+00  0.00000000e+00  0.00000000e+00
 0.00000000e+00  0.00000000e+00 -1.00000000e+01  0.00000000e+00
 1.00000000e+00 -0.00000000e+00 -0.00000000e+00 -4.37113883e-08
 0.00000000e+00  0.00000000e+00  0.00000000e+00  0.00000000e+00
 0.00000000e+00  0.00000000e+00  5.75471878e+00 -1.00000000e+00
 5.55726671e+00  0.00000000e+00  1.00000000e+00  0.00000000e+00
-1.68164849e-01]

```

### 1.0.3 3. Take Random Actions in the Environment

The cell below is just to give you an idea on how to make steps

Once this cell is executed, you will watch the agent's performance, if it selects an action at random with each time step. A window should pop up that allows you to observe the agent, as it moves through the environment.

```

[6]: env_info = env.reset(train_mode=False)[brain_name]      # reset the environment
    ↪
    states = env_info.vector_observations                    # get the current state
    ↪ (for each agent)
    scores = np.zeros(num_agents)                          # initialize the score
    ↪ (for each agent)
    while True:
        actions = np.random.randn(num_agents, action_size) # select an action (for
    ↪ each agent)
        actions = np.clip(actions, -1, 1)                  # all actions between -1
    ↪ and 1
        env_info = env.step(actions)[brain_name]           # send all actions to
    ↪ the environment
        next_states = env_info.vector_observations          # get next state (for
    ↪ each agent)
        rewards = env_info.rewards                         # get reward (for each
    ↪ agent)
        dones = env_info.local_done                        # see if episode finished
        scores += env_info.rewards                          # update the score (for
    ↪ each agent)
        states = next_states                                # roll over states to
    ↪ next time step
        if np.any(dones):                                  # exit loop if episode
    ↪ finished
            break
    print('Total score (averaged over agents) this episode: {}'.format(np.
    ↪ mean(scores)))

```

Total score (averaged over agents) this episode: 0.0

When finished, you can close the environment.

```
[9]: env.close()
```

#### 1.0.4 4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

## 2 My implementation

### 2.1 Actor Critic Algorithms

For infinite input space and infinite output space. Compared to Deep-Q that only allows a finite number of inputs, we have chosen Deep Deterministic Policy Gradient. ##### Deep Deterministic Policy Gradient Used for continuous action space, we have added noise to the process just like the authors did [Lillicrap et al, 2015](#).

#### 2.1.1 Use two separate models

- One outputs the desired action in the continuous space
- Other an action to produce Q-values

**Actor Critic Methods is about having two models:** Actor takes current environment state and determines the best action to take from there. Critic takes in state and action and return score of how good the action is.

**Here is the code for Actor and Critic Network:** Same architecture for both Agent and Critic with 2 fully connected layers of 400 and 300 units respectively, where values are normalized each batch. Activations are RELU on the first two layers for both Networks and then Tanh and no activation function respectively.

#### Hyperparameters

- `fc1_units=400`
- `fc2_units=300`
- `BUFFER_SIZE = 100000`
- `BATCH_SIZE = 64`
- `GAMMA = 0.9`

- $\text{TAU} = 1\text{e-}4$
- $\text{lr\_actor} = 1\text{e-}4$
- $\text{lr\_critic} = 1\text{e-}4$
- $\text{WEIGHT\_DECAY} = 0$
- $\text{LEARN\_EVERY} = 20$
- $\text{learning\_num} = 10$
- $\text{GRAD\_CLIPPING} = 1.0$
- $\text{ou\_sigma} = 0.2$
- $\text{ou\_theta} = 0.15$
- $\text{EPSILON} = 1.0$
- $\text{EPSILON\_DECAY} = 1\text{e-}6$

```
[11]: ! cat model.py
```

```
import numpy as np
import torch
import torch.nn as nn
import torch.nn.functional as F

class Actor(nn.Module):
    """Actor"""

    def __init__(self, state_size, action_size, seed, fc1_units=400,
fc2_units=300):
        """
        Using two fully connected layers with 400 and 300 units respectively.
        Not resetting parameters.
        """
        super(Actor, self).__init__()
        #use for
        self.seed = torch.manual_seed(seed)
        #for fully connected layer input
        self.fc1 = nn.Linear(state_size, fc1_units)
        #Applying batch normalization
        self.bn1 = nn.BatchNorm1d(fc1_units)
        #fully connected layers
        self.fc2 = nn.Linear(fc1_units, fc2_units)
        self.fc3 = nn.Linear(fc2_units, action_size)

    def forward(self, state):
```

```

        """Actor policy network to map states to actions, using relus and
tanh"""
        x = F.relu(self.bn1(self.fc1(state)))
        x = F.relu(self.fc2(x))
        return torch.tanh(self.fc3(x))

class Critic(nn.Module):
    """Critic"""
    def __init__(self, state_size, action_size, seed, fc1_units=400,
fc2_units=300):
        """Using same architecture as in Actor network. Two fully connected
layers of 400 and 300 units respectively.
        """
        super(Critic, self).__init__()
        self.seed = torch.manual_seed(seed)
        self.fc1 = nn.Linear(state_size, fc1_units)
        self.bn1 = nn.BatchNorm1d(fc1_units)
        #in forward function action will be concatenated according to DDGP
        self.fc2 = nn.Linear(fc1_units+action_size, fc2_units)
        self.fc3 = nn.Linear(fc2_units, 1)

    def forward(self, state, action):
        """Critic value network that maps (state,action) pairs to Q-values"""
        x = F.relu(self.bn1(self.fc1(state)))
        #to concatenate action
        x = torch.cat((x, action), dim=1)
        x = F.relu(self.fc2(x))
        return self.fc3(x)

```

```
[2]: ! cat agent.py
```

```

from collections import namedtuple, deque
import torch
import torch.nn.functional as F
import torch.optim as optim
import numpy as np
import random
from model import Actor, Critic
import copy

BUFFER_SIZE = 100000 # replay buffer
BATCH_SIZE = 64      # batch size: fixed batch per pass
GAMMA = 0.9          # discount factor
TAU = 1e-4           # for soft update: Not update at once but frequently
https://arxiv.org/pdf/1509.02971.pdf
lr_actor = 1e-4       # learning rate actor
lr_critic = 1e-4      # learning rate critic
WEIGHT_DECAY = 0      # L2 weight decay

```

```

LEARN_EVERY = 20          # learning timestep interval
learning_num = 10         # number of learning passes
GRAD_CLIPPING = 1.0      # gradient clipping

# Ornstein-Uhlenbeck: Stochastic stationary Gauss-Markov process
ou_sigma = 0.2
ou_theta = 0.15

EPSILON = 1.0
EPSILON_DECAY = 1e-6

#gpu if possible
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

class Agent():

    def __init__(self, state_size, action_size, seed=12):

        self.state_size = state_size
        self.action_size = action_size
        self.seed = random.seed(seed)

        self.epsilon = EPSILON

        #This the actor network. Imported from model.py
        self.actor_local = Actor(state_size, action_size, seed).to(device)
        self.actor_target = Actor(state_size, action_size, seed).to(device)
        self.actor_optimizer = optim.Adam(self.actor_local.parameters(),
lr=lr_actor)

        #We will test the critic to see how good is the action. Imported from
model.py
        self.critic_local = Critic(state_size, action_size, seed).to(device)
        self.critic_target = Critic(state_size, action_size, seed).to(device)
        self.critic_optimizer = optim.Adam(self.critic_local.parameters(),
lr=lr_critic, weight_decay=WEIGHT_DECAY)

        # From class UONoise
        self.noise = OUNoise(action_size, seed)

        # From class replay buffer
        self.memory = ReplayBuffer(action_size, BUFFER_SIZE, BATCH_SIZE, seed)

    def step(self, state, action, reward, next_state, done, timestep):
        """Save experience in replay buffer, and use sample from buffer to
learn"""

```

```

#from add function in Replay Buffer class: store experience
self.memory.add(state, action, reward, next_state, done)

# If having enough batch size
if len(self.memory) > BATCH_SIZE and timestep % LEARN_EVERY == 0:
    for _ in range(learning_num):
        #sample data from sample function Replay Buffer class
        experiences = self.memory.sample()
        #from learn function
        self.learn(experiences, GAMMA)

def act(self, state):
    """Returns actions for given state as per current policy"""
    #make state input
    state = torch.from_numpy(state).float().to(device)
    #in eval mode instead of training
    self.actor_local.eval()
    #do not save backprop but increase speed
    with torch.no_grad():
        action = self.actor_local(state).cpu().data.numpy()
    #train actor
    self.actor_local.train()

    #from sample noise class
    action += self.epsilon * self.noise.sample()

    return np.clip(action, -1, 1)

def reset(self):
    """ Reset from Noise to mean"""
    self.noise.reset()

def learn(self, experiences, gamma):
    """
    Experience is a tuple of states, actions, rewards, next_states, dones
    Update policy and value parameters using given batch of experience
    tuples.
    Q_targets = r +  $\gamma$  * critic_target(next_state, actor_target(next_state))
    """

    states, actions, rewards, next_states, dones = experiences

    #use target models
    actions_next = self.actor_target(next_states)
    #to feed critic we use actor actions
    Q_targets_next = self.critic_target(next_states, actions_next)
    # compute Q targets

```



```

    Q_targets = rewards + (gamma * Q_targets_next * (1 - dones))
    # critic loss
    Q_expected = self.critic_local(states, actions)
    critic_loss = F.mse_loss(Q_expected, Q_targets)
    # minimize critic the loss
    self.critic_optimizer.zero_grad()
    critic_loss.backward()
    # gradient clipping for critic
    if GRAD_CLIPPING > 0:
        torch.nn.utils.clip_grad_norm_(self.critic_local.parameters(),
GRAD_CLIPPING)
    #step updates the parameter
    self.critic_optimizer.step()

    # actor
    actions_pred = self.actor_local(states)
    actor_loss = -self.critic_local(states, actions_pred).mean()
    # minimize the loss
    self.actor_optimizer.zero_grad()
    actor_loss.backward()
    self.actor_optimizer.step()

    # update from soft_update function
    self.soft_update(self.critic_local, self.critic_target, TAU)
    self.soft_update(self.actor_local, self.actor_target, TAU)
    # update epsilon decay
    if EPSILON_DECAY > 0:
        self.epsilon -= EPSILON_DECAY
        self.noise.reset()

    def soft_update(self, local_model, target_model, tau):
        """Soft update model parameters.  = *_local + (1 - )*_target"""
        for target_param, local_param in zip(target_model.parameters(),
local_model.parameters()):
            target_param.data.copy_(tau*local_param.data +
(1.0-tau)*target_param.data)

class OUNoise:
    """Ornstein-Uhlenbeck"""

    def __init__(self, size, seed, mu=0., theta=ou_theta, sigma=ou_sigma):
        #array of ones
        self.mu = np.array(mu * size)
        self.theta = theta
        self.sigma = sigma
        self.seed = random.seed(seed)
        self.size = size
        self.reset()

```

```

def reset(self):
    """Return a shallow copy of x"""
    self.state = copy.copy(self.mu)

def sample(self):
    """Update internal state and return it as a noise sample"""
    state = self.state
    d_x = self.theta * (self.mu - state) + self.sigma *
np.random.standard_normal(self.size)
    self.state = state + d_x
    return self.state

class ReplayBuffer:
    """Store experience tuples."""

    def __init__(self, action_size, buffer_size, batch_size, seed):
        self.action_size = action_size
        #Deque (Doubly Ended Queue)
        self.memory = deque(maxlen=buffer_size)
        #size of each training
        self.batch_size = batch_size
        self.experience = namedtuple("Experience", field_names = ["state",
"action", "reward", "next_state", "done"])
        self.seed = random.seed(seed)

    def add(self, state, action, reward, next_state, done):
        """Add new experience to memory."""
        e = self.experience(state, action, reward, next_state, done)
        self.memory.append(e)

    def sample(self):
        """Random sample of batch from replay buffer"""
        experiences = random.sample(self.memory, k=self.batch_size)

        #vstack: Stack arrays in sequence vertically (row wise)
        states = torch.from_numpy(np.vstack([e.state for e in experiences if e
is not None])).float().to(device)
        actions = torch.from_numpy(np.vstack([e.action for e in experiences if e
is not None])).float().to(device)
        rewards = torch.from_numpy(np.vstack([e.reward for e in experiences if e
is not None])).float().to(device)
        next_states = torch.from_numpy(np.vstack([e.next_state for e in
experiences if e is not None])).float().to(device)
        dones = torch.from_numpy(np.vstack([e.done for e in experiences if e is
not None])).astype(np.uint8).float().to(device)

```

```
return (states, actions, rewards, next_states, dones)
```

```
def __len__(self):  
    """Size of internal memory"""  
    return len(self.memory)
```

```
[7]: import numpy as np  
import random  
import time  
import torch  
  
#for training performance  
import matplotlib.pyplot as plt  
%matplotlib inline  
  
from collections import deque  
from agent import Agent  
from unityagents import UnityEnvironment
```

```
[8]: def ddpq(actor_weights_name, critic_weights_name, num_episodes=1500,   
    ↪max_t=1000, print_every=25):  
    """Deep Deterministic Policy Gradient  
    """  
    #empty list for score storing  
    mean_scores = []  
    #empty list for moving average  
    moving_avgs = []  
    best_score = -np.inf  
    scores_window = deque(maxlen=100)  
    #iterate over number of episodes defined  
    for episode in range(1, num_episodes + 1):  
        #reset enviroment  
        env_info = env.reset(train_mode=True)[brain_name]  
        #state of observations  
        states = env_info.vector_observations  
        #set score to zero to number of agents  
        scores = np.zeros(num_agents)  
        #from agent script: reset  
        agent.reset()  
        #keep track of time  
        start_time = time.time()  
        for t in range(max_t):  
            #pick according to state  
            actions = agent.act(states)  
            #make decision according to actions  
            env_info = env.step(actions)[brain_name]  
            next_states = env_info.vector_observations
```

```

        #get rewards
        rewards = env_info.rewards
        #check for finished episode
        dones = env_info.local_done
        #learn (from agent fucntion)
        for state, action, reward, next_state, done in zip(states, actions,
→rewards, next_states, dones):
            agent.step(state, action, reward, next_state, done, t)
            states = next_states
            scores += rewards
            #when done break
            if np.any(dones):
                break
        end_time = time.time()
        duration = end_time - start_time
        #append mean score
        mean_scores.append(np.mean(scores))
        scores_window.append(mean_scores[-1])
        moving_avgs.append(np.mean(scores_window))

        if episode % print_every == 0:
            print("\rEpisode {} ({}s)\tMean: {:.1f}\tMoving Avg: {:.1f}"\
                .format(episode, round(duration), mean_scores[-1],
→moving_avgs[-1]))
            if moving_avgs[-1] >= 30.0:
                print("\nEnvironment solved in {:d} episodes.\tAverage score: {:.
→2f}"\
                    .format(episode, moving_avgs[-1]))
                torch.save(agent.actor_local.state_dict(), actor_weights_name)
                torch.save(agent.critic_local.state_dict(), critic_weights_name)
                break

        return(mean_scores, moving_avgs)

```

```
[9]: agent = Agent(state_size=state_size, action_size=action_size, seed=12)
```

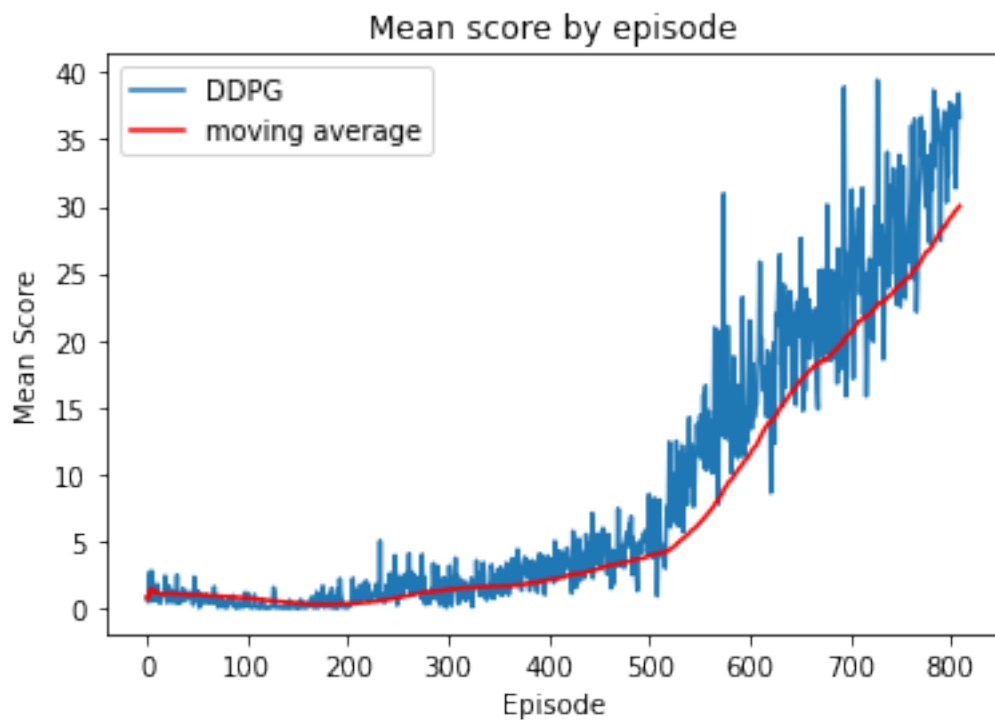
```
[10]: scores, avgs = ddpq(actor_weights_name='actor_single.pth',
→critic_weights_name='critic_single.pth', print_every=50)
```

Episode 50 (10s)	Mean: 0.5	Moving Avg: 1.0
Episode 100 (12s)	Mean: 0.1	Moving Avg: 0.7
Episode 150 (12s)	Mean: 0.4	Moving Avg: 0.3
Episode 200 (13s)	Mean: 0.0	Moving Avg: 0.3
Episode 250 (10s)	Mean: 0.2	Moving Avg: 0.7
Episode 300 (10s)	Mean: 2.0	Moving Avg: 1.3
Episode 350 (10s)	Mean: 1.1	Moving Avg: 1.6
Episode 400 (13s)	Mean: 3.9	Moving Avg: 2.1

Episode 450 (10s)	Mean: 5.5	Moving Avg: 2.9
Episode 500 (10s)	Mean: 3.8	Moving Avg: 3.8
Episode 550 (10s)	Mean: 12.2	Moving Avg: 6.2
Episode 600 (10s)	Mean: 17.9	Moving Avg: 11.3
Episode 650 (11s)	Mean: 23.0	Moving Avg: 16.8
Episode 700 (14s)	Mean: 26.5	Moving Avg: 20.5
Episode 750 (12s)	Mean: 33.8	Moving Avg: 24.2
Episode 800 (10s)	Mean: 37.7	Moving Avg: 29.1

Environment solved in 810 episodes.      Average score: 30.03

```
[14]: plt.plot(np.arange(len(scores)), scores, label='DDPG')
plt.plot(np.arange(len(scores)), avgs, c='r', label='moving average')
plt.title("Mean score by episode")
plt.ylabel('Mean Score')
plt.xlabel('Episode')
plt.legend(loc='upper left')
plt.show()
```



```
[15]: env.close()
```

## 2.2 Future Work:

- Even though Gamma value pretty high, we might be able to increase a little bit more.
- Decresing the learning interval or increasing the steps when learning might yield to faster results, also changing the Actor and Critic architecture.
- Finally by adding prioritized experience replay as [\(Hou & Zhang, 2017\)](#).